

Private SCT Auditing, Revisited

Lena Heimberger¹, Bas Westerbaan²

Graz University of Technology¹, Cloudflare²

November 20, 2025

Introduction

Introduction

- Certificates **bind** a domain name to a public key

Introduction

- Certificates **bind** a domain name to a public key
- add accountability to issuance process: public logging after **Maximum Merge Delay (MMD)**

Introduction

- Certificates **bind** a domain name to a public key
- add accountability to issuance process: public logging after **Maximum Merge Delay (MMD)**
- currently **six** logs that are widely trusted by browsers

Introduction

- Certificates **bind** a domain name to a public key
- add accountability to issuance process: public logging after **Maximum Merge Delay (MMD)**
- currently **six** logs that are widely trusted by browsers
- a certificate needs **two** valid SCTs(Signed Certificate Timestamps) for SCT-enforcing browsers

Introduction

- Certificates **bind** a domain name to a public key
- add accountability to issuance process: public logging after **Maximum Merge Delay (MMD)**
- currently **six** logs that are widely trusted by browsers
- a certificate needs **two** valid SCTs(Signed Certificate Timestamps) for SCT-enforcing browsers
- instead of trusting one CA: Trust two logs

Introduction

- Certificates **bind** a domain name to a public key
- add accountability to issuance process: public logging after **Maximum Merge Delay (MMD)**
- currently **six** logs that are widely trusted by browsers
- a certificate needs **two** valid SCTs(Signed Certificate Timestamps) for SCT-enforcing browsers
- instead of trusting one CA: Trust two logs
- more logs in the future: back to the start?

SCT Auditing

- SCTs need to be **checked for inclusion** in the log after the MMD

SCT Auditing

- SCTs need to be **checked for inclusion** in the log after the MMD
- checking directly is **problematic** for user privacy:

SCT Auditing

- SCTs need to be **checked for inclusion** in the log after the MMD
- checking directly is **problematic** for user privacy: leaks browser history

SCT Auditing

- SCTs need to be **checked for inclusion** in the log after the MMD
- checking directly is **problematic** for user privacy: leaks browser history
- very limited auditing by differentially private lookups

SCT Auditing

- SCTs need to be **checked for inclusion** in the log after the MMD
- checking directly is **problematic** for user privacy: leaks browser history
- very limited auditing by differentially private lookups
- Can we do better (e.g., PIR?)

SCT Auditing

- SCTs need to be **checked for inclusion** in the log after the MMD
- checking directly is **problematic** for user privacy: leaks browser history
- very limited auditing by differentially private lookups
- Can we do better (e.g., PIR?)
Earlier analysis suggests this is impractical

This Paper Presentation

1. SCT auditing and **slow issuance certificates**: post-quantum certificates, alternatives to SCTs, and rough estimates

This Paper Presentation

1. SCT auditing and **slow issuance certificates**: post-quantum certificates, alternatives to SCTs, and rough estimates
2. revisiting SCT auditing with **concrete requirements** for PIR, and how the state of the art fulfills our requirements

This Paper Presentation

1. SCT auditing and **slow issuance certificates**: post-quantum certificates, alternatives to SCTs, and rough estimates
2. revisiting SCT auditing with **concrete requirements** for PIR, and how the state of the art fulfills our requirements
3. Benchmarks with ideas on database management to get an **efficient PIR lookup**

This Paper Presentation

1. SCT auditing and **slow issuance certificates**: post-quantum certificates, alternatives to SCTs, and rough estimates
2. revisiting SCT auditing with **concrete requirements** for PIR, and how the state of the art fulfills our requirements
3. Benchmarks with ideas on database management to get an **efficient PIR lookup**
4. Batching audits at little to no cost to scale to audits **under 100 kB per day**

Preliminaries

Certificate Transparency

- 2011: DigiNotar wrongfully issued over 500 certificates

Certificate Transparency

- 2011: DigiNotar wrongfully issued over 500 certificates led to CA being distrusted!

Certificate Transparency

- 2011: DigiNotar wrongfully issued over 500 certificates led to CA being distrusted!
- CT requires all certificates to be **publicly logged**

Certificate Transparency

- 2011: DigiNotar wrongfully issued over 500 certificates led to CA being distrusted!
- CT requires all certificates to be publicly logged
- after logging, a log returns a Signed Certificate Timestamp (SCT)

Certificate Transparency

- 2011: DigiNotar wrongfully issued over 500 certificates led to CA being distrusted!
- CT requires all certificates to be publicly logged
- after logging, a log returns a Signed Certificate Timestamp (SCT)
- Firefox, Safari, Chrome, and Brave require at least two SCTs from distinct, trusted logs to validate a certificate.

Certificate Transparency

- 2011: DigiNotar wrongfully issued over 500 certificates led to CA being distrusted!
- CT requires all certificates to be **publicly logged**
- after logging, a log returns a **Signed Certificate Timestamp** (SCT)
- Firefox, Safari, Chrome, and Brave require at least two SCTs from distinct, trusted logs to validate a certificate.
- the log datastructure enables the computation of two proofs:

Certificate Transparency

- 2011: DigiNotar wrongfully issued over 500 certificates led to CA being distrusted!
- CT requires all certificates to be publicly logged
- after logging, a log returns a Signed Certificate Timestamp (SCT)
- Firefox, Safari, Chrome, and Brave require at least two SCTs from distinct, trusted logs to validate a certificate.
- the log datastructure enables the computation of two proofs: An inclusion proof that proves a specific node is included in the log,

Certificate Transparency

- 2011: DigiNotar wrongfully issued over 500 certificates led to CA being distrusted!
- CT requires all certificates to be **publicly logged**
- after logging, a log returns a **Signed Certificate Timestamp (SCT)**
- Firefox, Safari, Chrome, and Brave require at least two SCTs from distinct, trusted logs to validate a certificate.
- the log datastructure enables the computation of two proofs: An **inclusion** proof that proves a specific node is included in the log, and a **consistency** proof proving there are no split-view attacks

Private SCT auditing

- Controlling two CT logs is sufficient to serve a malicious certificate without immediate detection by the browser

Private SCT auditing

- Controlling two CT logs is sufficient to serve a malicious certificate without immediate detection by the browser
- Two phases in SCT auditing:

Private SCT auditing

- Controlling two CT logs is sufficient to serve a malicious certificate without immediate detection by the browser
- Two phases in SCT auditing:
 1. **reporting** phase where an SCT is collected

Private SCT auditing

- Controlling two CT logs is sufficient to serve a malicious certificate without immediate detection by the browser
- Two phases in SCT auditing:
 1. **reporting** phase where an SCT is collected
 2. **auditing** phase where an SCT is checked for inclusion

Private SCT auditing

- Controlling two CT logs is sufficient to serve a malicious certificate without immediate detection by the browser
- Two phases in SCT auditing:
 1. **reporting** phase where an SCT is collected
 2. **auditing** phase where an SCT is checked for inclusion
- Public logging of collected SCTs is sufficient!

Private SCT auditing

- Controlling two CT logs is sufficient to serve a malicious certificate without immediate detection by the browser
- Two phases in SCT auditing:
 1. **reporting** phase where an SCT is collected
 2. **auditing** phase where an SCT is checked for inclusion
- Public logging of collected SCTs is sufficient!
- domain owners can **monitor** logs to check for misbehaviour

SCT Auditing in Practice: Google Chrome

- Chrome **audits** SCTs

SCT Auditing in Practice: Google Chrome

- Chrome audits SCTs
- SCTs from popular websites are stored in the browser

SCT Auditing in Practice: Google Chrome

- Chrome audits SCTs
- SCTs from popular websites are stored in the browser
- PwnedPasswords-like approach: query hash prefix, get hash postfixes.

SCT Auditing in Practice: Google Chrome

- Chrome audits SCTs
- SCTs from popular websites are stored in the browser
- PwnedPasswords-like approach: query hash prefix, get hash postfixes.
- possible network level attack-selectively target communication and drop it

SCT Auditing in Practice: Google Chrome

- Chrome audits SCTs
- SCTs from popular websites are stored in the browser
- PwnedPasswords-like approach: query hash prefix, get hash postfixes.
- possible network level attack-selectively target communication and drop it
 - persist-and-retry mechanism

Private Information Retrieval

- Privately retrieve an item from database

Private Information Retrieval

- Privately retrieve an item from database
- for SCT auditing: Trillian did not distinguish between sequencing and including an entry until last year

Private Information Retrieval

- Privately retrieve an item from database
- for SCT auditing: Trillian did not distinguish between sequencing and including an entry until last year
lookup inclusion **by hash!**

Private Information Retrieval

- Privately retrieve an item from database
- for SCT auditing: Trillian did not distinguish between sequencing and including an entry until last year
lookup inclusion **by hash!**
- hash certificates into Bloom filter, get single-bit response indicating if record exists

Private Information Retrieval

- Privately retrieve an item from database
- for SCT auditing: Trillian did not distinguish between sequencing and including an entry until last year
lookup inclusion **by hash!**
- hash certificates into Bloom filter, get single-bit response indicating if record exists
recompute Bloom filter on the server-side when a new certificate is added.

Private Information Retrieval

- Privately retrieve an item from database
- for SCT auditing: Trillian did not distinguish between sequencing and including an entry until last year
lookup inclusion **by hash!**
- hash certificates into Bloom filter, get single-bit response indicating if record exists
recompute Bloom filter on the server-side when a new certificate is added.
- PIR protocols improved: more **efficient, single-server** protocols

Private Information Retrieval

- Privately retrieve an item from database
- for SCT auditing: Trillian did not distinguish between sequencing and including an entry until last year
lookup inclusion **by hash!**
- hash certificates into Bloom filter, get single-bit response indicating if record exists
recompute Bloom filter on the server-side when a new certificate is added.
- PIR protocols improved: more **efficient, single-server** protocols
- Can we do a lookup **by sequence number** and get the hash as a response?

Slow Issuance Certificates

Slow-issuance certificates

- elliptic curve signatures are 65 bytes

Slow-issuance certificates

- elliptic curve signatures are 65 bytes
- ML-DSA-44 signatures are 2.4 kB

Slow-issuance certificates

- elliptic curve signatures are 65 bytes
- ML-DSA-44 signatures are 2.4 kB
- Adding ML-DSA-44 to a TLS handshake will **more than double** the size of the handshake

Slow-issuance certificates

- elliptic curve signatures are 65 bytes
- ML-DSA-44 signatures are 2.4 kB
- Adding ML-DSA-44 to a TLS handshake will **more than double** the size of the handshake
→ switch to non-X.509 PKI

Slow-issuance certificates

- elliptic curve signatures are 65 bytes
- ML-DSA-44 signatures are 2.4 kB
- Adding ML-DSA-44 to a TLS handshake will **more than double** the size of the handshake
→ switch to non-X.509 PKI
- Embedding a proof-of-inclusion in the certificate (instead of an SCT) would remove the need for auditing

Slow-issuance certificates

- elliptic curve signatures are 65 bytes
- ML-DSA-44 signatures are 2.4 kB
- Adding ML-DSA-44 to a TLS handshake will **more than double** the size of the handshake
→ switch to non-X.509 PKI
- Embedding a proof-of-inclusion in the certificate (instead of an SCT) would remove the need for auditing
- Chrome says ML-DSA-44 is **undeployable**

The problem of SCTs and post-quantum certificates

- Right now, certificates can be used **immediately** after they are issued: this is why SCT audits are necessary

The problem of SCTs and post-quantum certificates

- Right now, certificates can be used **immediately** after they are issued: this is why SCT audits are necessary
- giving up immediate issuance eliminates the need for SCTs and SCT auditing

The problem of SCTs and post-quantum certificates

- Right now, certificates can be used **immediately** after they are issued: this is why SCT audits are necessary
- giving up immediate issuance eliminates the need for SCTs and SCT auditing
- use **proof of inclusion**, certificate is usable after new tree heads propagate

The problem of SCTs and post-quantum certificates

- Right now, certificates can be used **immediately** after they are issued: this is why SCT audits are necessary
- giving up immediate issuance eliminates the need for SCTs and SCT auditing
- use **proof of inclusion**, certificate is usable after new tree heads propagate
- When is a new certificate needed immediately?

Scenarios for immediate issuance

1. **Brand New Domains** that were just created. Low traffic, but once for each domain.

Scenarios for immediate issuance

1. **Brand New Domains** that were just created. Low traffic, but once for each domain.
2. **Validity Gaps** from configuration errors. Hopefully disappears over time.

Scenarios for immediate issuance

1. **Brand New Domains** that were just created. Low traffic, but **once for each domain**.
2. **Validity Gaps** from configuration errors. Hopefully **disappears** over time.
3. **Urgent Re-Issuance** when a domain is **moved on short notice**, e.g., because of a DDoS attack.

Client Support

1. Supporting **with** a reliable update mechanism, rarely requires fallbacks

Client Support

1. Supporting **with** a reliable update mechanism, rarely requires fallbacks
2. Supporting **without** getting reliable updates, falls back more often when becoming **stale**

Client Support

1. Supporting **with** a reliable update mechanism, rarely requires fallbacks
2. Supporting **without** getting reliable updates, falls back more often when becoming **stale**
3. No support, requires the fallback mechanism **every time**

Estimating a Slow-Issuance Ecosystem

- ecosystem does not exist, but we can try to estimate server behaviour

Estimating a Slow-Issuance Ecosystem

- ecosystem does not exist, but we can try to estimate server behaviour
- assuming X.509 certificates are replaced with slow-issuance certificates, we have two datasets:

Estimating a Slow-Issuance Ecosystem

- ecosystem does not exist, but we can try to estimate server behaviour
- assuming X.509 certificates are replaced with slow-issuance certificates, we have two datasets:
 - random domains with 616 142 certificates

Estimating a Slow-Issuance Ecosystem

- ecosystem does not exist, but we can try to estimate server behaviour
- assuming X.509 certificates are replaced with slow-issuance certificates, we have two datasets:
 - **random domains** with 616 142 certificates corresponding to 123 733 unique domains

Estimating a Slow-Issuance Ecosystem

- ecosystem does not exist, but we can try to estimate server behaviour
- assuming X.509 certificates are replaced with slow-issuance certificates, we have two datasets:
 - **random domains** with 616 142 certificates corresponding to 123 733 unique domains
 - **top domains** with 36 375 certificates from 75 out of the top 100 domains

Estimating a Slow-Issuance Ecosystem

- ecosystem does not exist, but we can try to estimate server behaviour
- assuming X.509 certificates are replaced with slow-issuance certificates, we have two datasets:
 - **random domains** with 616 142 certificates corresponding to 123 733 unique domains
 - **top domains** with 36 375 certificates from 75 out of the top 100 domains
- Brand new domains:

Estimating a Slow-Issuance Ecosystem

- ecosystem does not exist, but we can try to estimate server behaviour
- assuming X.509 certificates are replaced with slow-issuance certificates, we have two datasets:
 - **random domains** with 616 142 certificates corresponding to 123 733 unique domains
 - **top domains** with 36 375 certificates from 75 out of the top 100 domains
- Brand new domains: $\approx 95\,000$ new domains

Estimating Validity Gaps

- slow embedding crucially relies on correct updating mechanisms

Estimating Validity Gaps

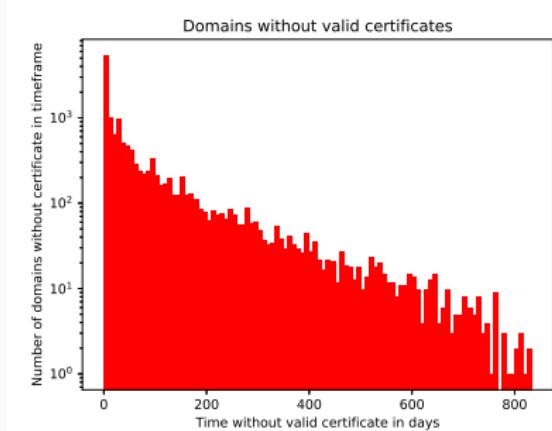
- slow embedding crucially relies on correct updating mechanisms
- Our statistics indicate that most domains without a valid certificate are renewed within a day

Estimating Validity Gaps

- slow embedding crucially relies on correct updating mechanisms
- Our statistics indicate that most domains without a valid certificate are renewed within a day
- we hope wrong configurations are eliminated over time

Estimating Validity Gaps

- slow embedding crucially relies on correct updating mechanisms
- Our statistics indicate that most domains without a valid certificate are renewed within a day
- we hope wrong configurations are eliminated over time
- 100 000 domains per day



Estimating Domain Moves

- Certificates have clear renewal patterns

Estimating Domain Moves

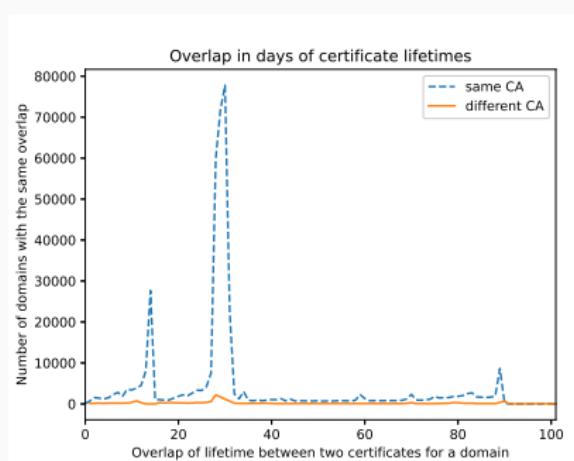
- Certificates have clear renewal patterns
- new CA + renewal outside of pattern + no other valid certificate present: counted as triggering a fallback

Estimating Domain Moves

- Certificates have clear renewal patterns
- new CA + renewal outside of pattern + no other valid certificate present: counted as triggering a fallback
- 100 000 certificates per day

Estimating Domain Moves

- Certificates have clear renewal patterns
- new CA + renewal outside of pattern + no other valid certificate present: counted as triggering a fallback
- 100 000 certificates per day



Towards Fully Private Auditing



\mathcal{SCT} obtained!

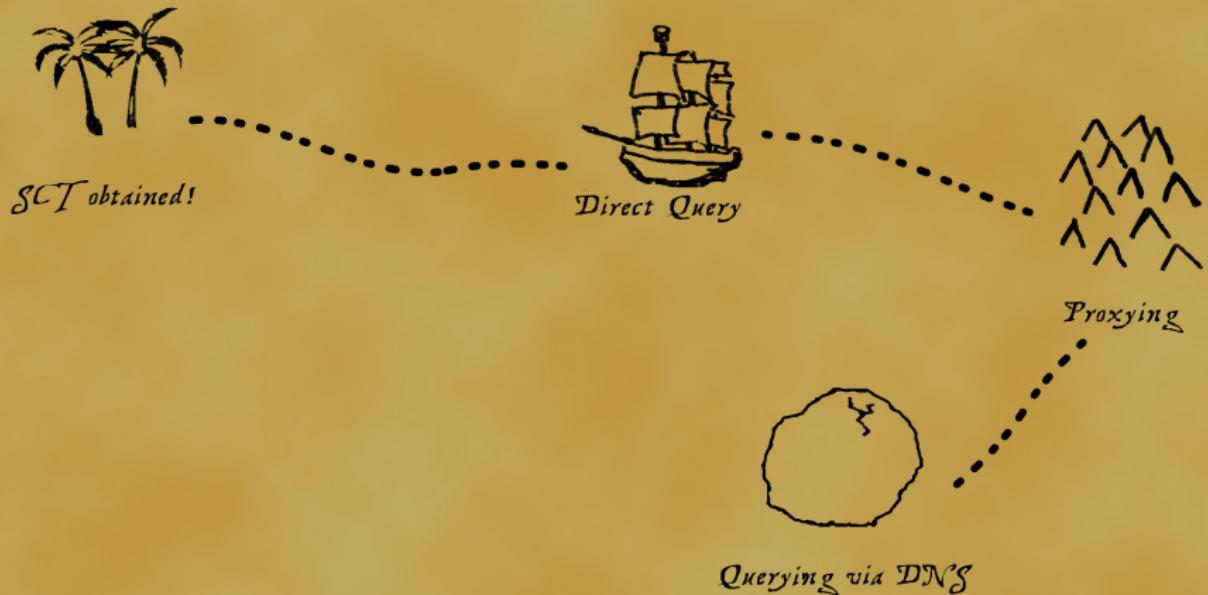


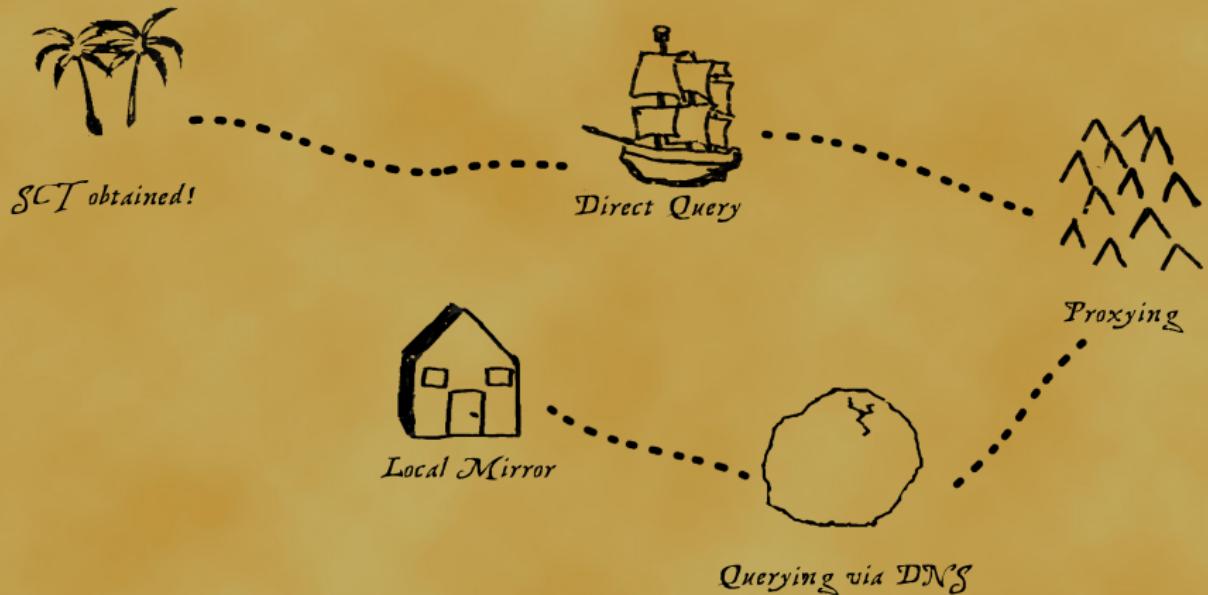
SCT obtained!

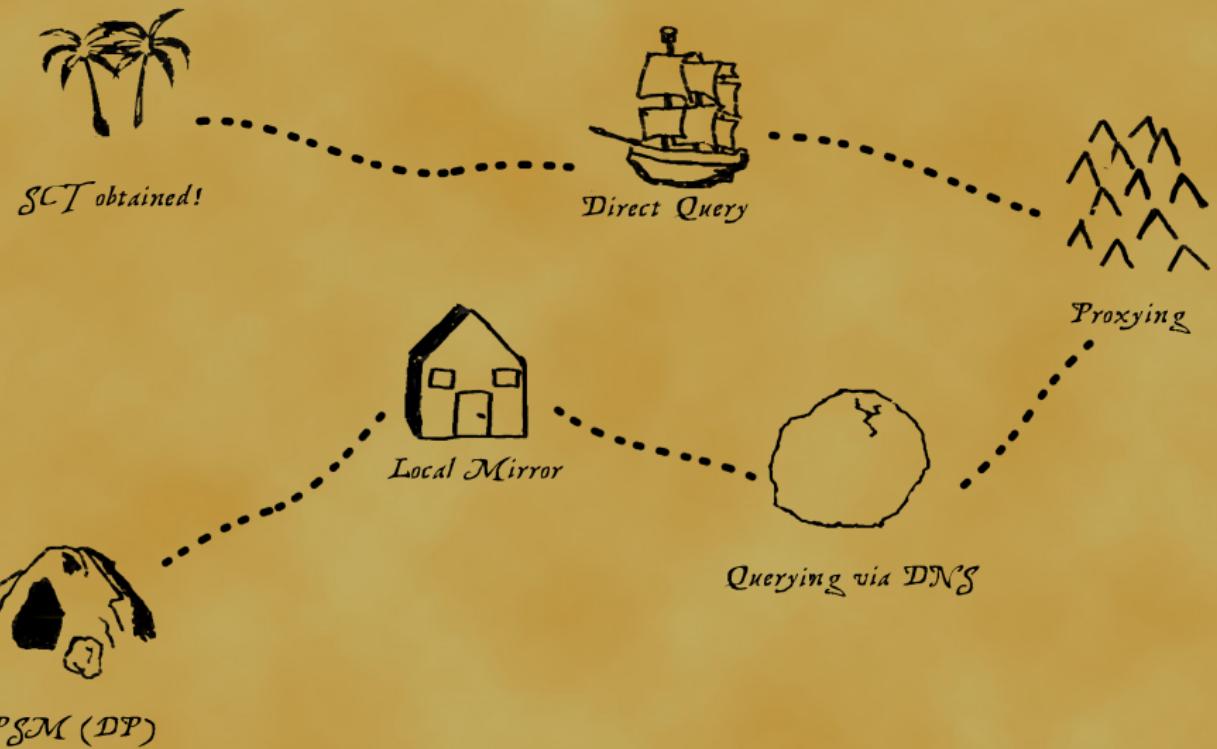


Direct Query











Constraints for SCT auditing with PIR

- **Lookup Requirements** lookup certificate **hash by sequence** number with **trusted third party**, e.g. browser auditing infrastructure.

Constraints for SCT auditing with PIR

- **Lookup Requirements** lookup certificate **hash by sequence** number with **trusted third party**, e.g. browser auditing infrastructure.
- **Database Preprocessing** minimize **per-user**, user-independent is feasible

Constraints for SCT auditing with PIR

- **Lookup Requirements** lookup certificate **hash by sequence** number with **trusted third party**, e.g. browser auditing infrastructure.
- **Database Preprocessing** minimize **per-user**, user-independent is feasible
- **Batching** Audits are ideally performed **immediately**, but batching or simplified batching are within scope

Constraints for SCT auditing with PIR

- **Lookup Requirements** lookup certificate **hash by sequence** number with **trusted third party**, e.g. browser auditing infrastructure.
- **Database Preprocessing** minimize **per-user**, user-independent is feasible
- **Batching** Audits are ideally performed **immediately**, but batching or simplified batching are within scope
- **Communication** $\approx 1\,000$ TLS connections/day, each 2.5-5 kB handshake.

Constraints for SCT auditing with PIR

- **Lookup Requirements** lookup certificate **hash by sequence** number with **trusted third party**, e.g. browser auditing infrastructure.
- **Database Preprocessing** minimize **per-user**, user-independent is feasible
- **Batching** Audits are ideally performed **immediately**, but batching or simplified batching are within scope
- **Communication** $\approx 1\,000$ TLS connections/day, each 2.5-5 kB handshake. A fractional **daily increase** of 100 kB is also tolerable

Constraints for SCT auditing with PIR

- **Lookup Requirements** lookup certificate **hash by sequence** number with **trusted third party**, e.g. browser auditing infrastructure.
- **Database Preprocessing** minimize **per-user**, user-independent is feasible
- **Batching** Audits are ideally performed **immediately**, but batching or simplified batching are within scope
- **Communication** $\approx 1\,000$ TLS connections/day, each 2.5-5 kB handshake. A fractional **daily increase** of 100 kB is also tolerable
- **Compute** One core-second per user per day

Constraints for SCT auditing with PIR

- **Lookup Requirements** lookup certificate **hash by sequence** number with **trusted third party**, e.g. browser auditing infrastructure.
- **Database Preprocessing** minimize **per-user**, user-independent is feasible
- **Batching** Audits are ideally performed **immediately**, but batching or simplified batching are within scope
- **Communication** $\approx 1\,000$ TLS connections/day, each 2.5-5 kB handshake. A fractional **daily increase** of 100 kB is also tolerable
- **Compute** One core-second per user per day \rightarrow a thousand 64-core servers

Constraints for SCT auditing with PIR

- **Lookup Requirements** lookup certificate **hash by sequence** number with **trusted third party**, e.g. browser auditing infrastructure.
- **Database Preprocessing** minimize **per-user**, user-independent is feasible
- **Batching** Audits are ideally performed **immediately**, but batching or simplified batching are within scope
- **Communication** $\approx 1\,000$ TLS connections/day, each 2.5-5 kB handshake. A fractional **daily increase** of 100 kB is also tolerable
- **Compute** One core-second per user per day \rightarrow a thousand 64-core servers
minimize server compute overhead (e.g. updating Bloom filters)

Constraints for SCT auditing with PIR

- **Lookup Requirements** lookup certificate **hash by sequence** number with **trusted third party**, e.g. browser auditing infrastructure.
- **Database Preprocessing** minimize **per-user**, user-independent is feasible
- **Batching** Audits are ideally performed **immediately**, but batching or simplified batching are within scope
- **Communication** $\approx 1\,000$ TLS connections/day, each 2.5-5 kB handshake. A fractional **daily increase** of 100 kB is also tolerable
- **Compute** One core-second per user per day \rightarrow a thousand 64-core servers
minimize server compute overhead (e.g. updating Bloom filters)
- **Reusing key material** Reusing public keys is acceptable for private audits

Constraints for SCT auditing with PIR

- **Lookup Requirements** lookup certificate **hash by sequence** number with **trusted third party**, e.g. browser auditing infrastructure.
- **Database Preprocessing** minimize **per-user**, user-independent is feasible
- **Batching** Audits are ideally performed **immediately**, but batching or simplified batching are within scope
- **Communication** $\approx 1\,000$ TLS connections/day, each 2.5-5 kB handshake. A fractional **daily increase** of 100 kB is also tolerable
- **Compute** One core-second per user per day \rightarrow a thousand 64-core servers
minimize server compute overhead (e.g. updating Bloom filters)
- **Reusing key material** Reusing public keys is acceptable for private audits - allows linking queries -

Constraints for SCT auditing with PIR

- **Lookup Requirements** lookup certificate **hash by sequence** number with **trusted third party**, e.g. browser auditing infrastructure.
- **Database Preprocessing** minimize **per-user**, user-independent is feasible
- **Batching** Audits are ideally performed **immediately**, but batching or simplified batching are within scope
- **Communication** $\approx 1\,000$ TLS connections/day, each 2.5-5 kB handshake. A fractional **daily increase** of 100 kB is also tolerable
- **Compute** One core-second per user per day \rightarrow a thousand 64-core servers
minimize server compute overhead (e.g. updating Bloom filters)
- **Reusing key material** Reusing public keys is acceptable for private audits - allows linking queries - add randomized delays

PIR from Fully Homomorphic Encryption

- SealPIR encrypts the record index as a polynomial x^i , privately expands it into a base vector and performs oblivious multiplication with the database

PIR: State of the Art

PIR from Fully Homomorphic Encryption

- SealPIR encrypts the record index as a polynomial x^i , privately expands it into a base vector and performs oblivious multiplication with the database
- Queries can get smaller with recursion at the cost of larger public parameters and per-client storage

PIR from Fully Homomorphic Encryption

- SealPIR encrypts the record index as a polynomial x^i , privately expands it into a base vector and performs oblivious multiplication with the database
- Queries can get smaller with recursion at the cost of larger public parameters and per-client storage

PIR from Learning-With-Errors

- LWE-based schemes send a base vector indicating the index

PIR: State of the Art

PIR from Fully Homomorphic Encryption

- SealPIR encrypts the record index as a polynomial x^i , privately expands it into a base vector and performs oblivious multiplication with the database
- Queries can get smaller with recursion at the cost of larger public parameters and per-client storage

PIR from Learning-With-Errors

- LWE-based schemes send a base vector indicating the index
- LWE is faster for computation, RLWE has smaller elements

PIR: State of the Art

PIR from Fully Homomorphic Encryption

- SealPIR encrypts the record index as a polynomial x^i , privately expands it into a base vector and performs oblivious multiplication with the database
- Queries can get smaller with recursion at the cost of larger public parameters and per-client storage

PIR from Learning-With-Errors

- LWE-based schemes send a base vector indicating the index
- LWE is faster for computation, RLWE has smaller elements
- HintlessPIR and TiptoePIR eliminate the hint by bootstrapping: outer RLWE element with inner LWE element

PIR: State of the Art

PIR from Fully Homomorphic Encryption

- SealPIR encrypts the record index as a polynomial x^i , privately expands it into a base vector and performs oblivious multiplication with the database
- Queries can get smaller with recursion at the cost of larger public parameters and per-client storage

PIR from Learning-With-Errors

- LWE-based schemes send a base vector indicating the index
- LWE is faster for computation, RLWE has smaller elements
- HintlessPIR and TiptoePIR eliminate the hint by bootstrapping: outer RLWE element with inner LWE element
- YPIR uses an algebraic transformation between RLWE and LWE, tiny responses

Downsizing the Database

Downsizing the Database

- PIR schemes perform a lot better with smaller databases

Downsizing the Database

- PIR schemes perform a lot better with smaller databases
- SCT Merkle trees grow steadily in size

Downsizing the Database

- PIR schemes perform a lot better with smaller databases
- SCT Merkle trees grow steadily in size
- **shard** database for better performance

Sharding by Structure: Tiling

- SCT logs group entries into tiles

Sharding by Structure: Tiling

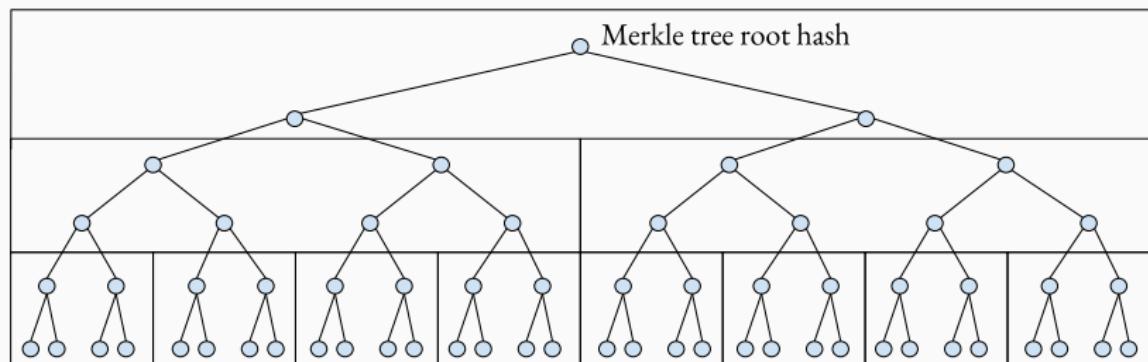
- SCT logs group entries into **tiles**
- check whether tile is in the tree

Sharding by Structure: Tiling

- SCT logs group entries into **tiles**
- check whether tile is in the tree
- remaining tile has to be included in the certificate or database record

Sharding by Structure: Tiling

- SCT logs group entries into **tiles**
- check whether tile is in the tree
- remaining tile has to be included in the certificate or database record
- $2^{15.5}$ to $2^{23.5}$ database records (2 levels up, one level up)



Each rectangle represents a tile with a height of 2, containing up to 6 hashes

Sharding by Epoch:

- CT logs grow constantly, but we only want to audit valid certificates!

Sharding by Epoch:

- CT logs grow constantly, but we only want to audit valid certificates!
- previously discussed as **STH discipline**

Sharding by Epoch:

- CT logs grow constantly, but we only want to audit valid certificates!
- previously discussed as **STH discipline**
- group certificates and publish new **temporal treeheads** periodically, e.g. every second

Sharding by Epoch:

- CT logs grow constantly, but we only want to audit valid certificates!
- previously discussed as **STH discipline**
- group certificates and publish new **temporal treeheads** periodically, e.g. every second
- audit temporal tree head

Sharding by Epoch:

- CT logs grow constantly, but we only want to audit valid certificates!
- previously discussed as **STH discipline**
- group certificates and publish new **temporal treeheads** periodically, e.g. every second
- audit temporal tree head
- further reduction: group certificates together again after a longer period, e.g., a day.

Sharding by Epoch:

- CT logs grow constantly, but we only want to audit valid certificates!
- previously discussed as **STH discipline**
- group certificates and publish new **temporal treeheads** periodically, e.g. every second
- audit temporal tree head
- further reduction: group certificates together again after a longer period, e.g., a day.
- especially interesting for the **fast-issuance fallback** case

Benchmarks

| Database parameters | Scheme | Request Size | Response Size |
|-------------------------------------|--|---|---|
| $ DB = 2^{31.5}$ $ r = 1$ | plain YPIR | 735 kB | 12 kB |
| $ DB = 2^{15.5}$ $ r = 256$ | SealPIR $d = 1$ SealPIR $d = 2$ Hintless PIR YPIR-SimplePIR | 46 kB 93 kB 370 kB 932 kB | 46 kB 185 kB 5.9 MB 12 kB |
| $ DB = 2^{15.5}$ $ r = 131072$ | SealPIR $d = 1$ SealPIR $d = 2$ Hintless PIR YPIR-SimplePIR | - - - 932 kB | - - - 61 kB |
| $ DB = 2^{17}$ $ r = 256$ | SealPIR $d = 1$ SealPIR $d = 2$ Hintless PIR YPIR-SimplePIR | 46 kB 93 kB 371 kB 1.4 MB | 46 kB 186 kB 6.0 MB 12 kB |

Batch Auditing

Batch Auditing

- auditing 1 000 certificates per day at 46 kB is prohibitive

Batch Auditing

- auditing 1 000 certificates per day at 46 kB is prohibitive
- combine STH discipline with slow-issuance certificates: a single audit per day

Batch Auditing

- auditing 1 000 certificates per day at 46 kB is prohibitive
- combine STH discipline with slow-issuance certificates: a single audit per day
- amortize PIR cost by packing multiple queries into one per day

Batch Auditing

- auditing 1 000 certificates per day at 46 kB is prohibitive
- combine STH discipline with slow-issuance certificates: a single audit per day
- amortize PIR cost by packing multiple queries into one per day
- sum of hashes is enough

Private Retrieval of Sums of Hashes

- LinPIR describes querying arbitrary linear combinations
 $\sum_i a_i \text{DB}_i$

Private Retrieval of Sums of Hashes

- LinPIR describes querying arbitrary linear combinations $\sum_i a_i \text{DB}_i$
- Batching with LWE: Encode multiple indices in the query vector

Private Retrieval of Sums of Hashes

- LinPIR describes querying arbitrary linear combinations $\sum_i a_i \text{DB}_i$
- Batching with LWE: Encode multiple indices in the query vector
- Batching with FHE: Encode multiple indices in the polynomial

Private Retrieval of Sums of Hashes

- LinPIR describes querying arbitrary linear combinations $\sum_i a_i \text{DB}_i$
- Batching with LWE: Encode multiple indices in the query vector
- Batching with FHE: Encode multiple indices in the polynomial $x^i + x^j$

Private Retrieval of Sums of Hashes

- LinPIR describes querying arbitrary linear combinations
 $\sum_i a_i \text{DB}_i$
- Batching with LWE: Encode multiple indices in the query vector
- Batching with FHE: Encode multiple indices in the polynomial $x^i + x^j$
expansion is linear

Security of retrieving Sums of Hashes

- Without batching: n -bit hash needs to be collision resistant

Security of retrieving Sums of Hashes

- Without batching: n -bit hash needs to be collision resistant
- retrieving sums: AdHASH map needs to be collision resistant

Security of retrieving Sums of Hashes

- Without batching: n -bit hash needs to be collision resistant
- retrieving sums: AdHASH map needs to be collision resistant(against published hashes)

Security of retrieving Sums of Hashes

- Without batching: n -bit hash needs to be collision resistant
- retrieving sums: AdHASH map needs to be collision resistant(against published hashes)
- Wagner's algorithm: Collisions in $O(k2^{\frac{n}{2+\lg k}+1})$ for k elements

Security of retrieving Sums of Hashes

- Without batching: n -bit hash needs to be collision resistant
- retrieving sums: AdHASH map needs to be collision resistant(against published hashes)
- Wagner's algorithm: Collisions in $O(k2^{\frac{n}{2+\lg k}+1})$ for k elements
- unclear if the attack is optimal: lower bound is $\Omega(2^{\frac{n}{2k}})$

Retrieving signed sums of hashes

- Solution: leverage algebraic structure to introduce noise

Retrieving signed sums of hashes

- Solution: leverage algebraic structure to introduce noise
- query **signed sum** or **multiples of elements**

Retrieving signed sums of hashes

- Solution: leverage algebraic structure to introduce noise
- query **signed sum** or **multiples of elements**
- e.g. $x^i + x^j + x^k \rightarrow -2x^i + 5x^j + x^k$

Conclusion

Conclusion

- SCT auditing with PIR is possible, thanks to sequence numbers and better PIR proposals

Conclusion

- SCT auditing with PIR is possible, thanks to sequence numbers and better PIR proposals
- Batching or slow-issuance certificates keep it under 100 kB/day

Conclusion

- SCT auditing with PIR is possible, thanks to sequence numbers and better PIR proposals
- Batching or slow-issuance certificates keep it under 100 kB/day
- Future Work
 - Estimate the slow-issuance ecosystem by slowly rolling it out

Conclusion

- SCT auditing with PIR is possible, thanks to sequence numbers and better PIR proposals
- Batching or slow-issuance certificates keep it under 100 kB/day
- Future Work
 - Estimate the slow-issuance ecosystem by slowly rolling it out
 - SCT-friendly PIR: probably some design space

Conclusion

- SCT auditing with PIR is possible, thanks to sequence numbers and better PIR proposals
- Batching or slow-issuance certificates keep it under 100 kB/day
- Future Work
 - Estimate the slow-issuance ecosystem by slowly rolling it out
 - SCT-friendly PIR: probably some design space
 - Practical challenges for deployment:

Conclusion

- SCT auditing with PIR is possible, thanks to sequence numbers and better PIR proposals
- Batching or slow-issuance certificates keep it under 100 kB/day
- Future Work
 - Estimate the slow-issuance ecosystem by slowly rolling it out
 - SCT-friendly PIR: probably some design space
 - Practical challenges for deployment:
 - blocked traffic/captive portals

Conclusion

- SCT auditing with PIR is possible, thanks to sequence numbers and better PIR proposals
- Batching or slow-issuance certificates keep it under 100 kB/day
- Future Work
 - Estimate the slow-issuance ecosystem by slowly rolling it out
 - SCT-friendly PIR: probably some design space
 - Practical challenges for deployment:
 - blocked traffic/captive portals
 - reporting open problem: **actionable and concrete evidence of log's misbehavior**